


Fewer failed pairs with each LLM configuration

Same 200 pairs · lower is better · paired adjusted $p < 0.001$ versus either local rule method



Exploratory matched comparisons. Reports share 16 underlying patterns; treating every pair as independent can overstate certainty.
Pattern-adjusted $p = 0.156$ versus earlier rules and 0.375 versus refined rules. The refined comparison was added after the original analysis.